

## Homework Assignment 2: Numpy Array and DataFrame

## 1. Define a function to analyze a numpy array

- Assume we have an array  $X$  which contains term frequency of each document. In this array, each row presents a document, each column denotes a word, and each value, say  $x_{i,j}$ , denotes the frequency of the word  $j$  in document  $i$ . Therefore, if there are  $m$  documents,  $n$  words,  $X$  has a shape of  $(m, n)$ .

Define a function named `analyze_tf` which:

- take  $X$  as an input.
  - calculate the document frequency  $df_j$  for word  $j$ , e.g. how many documents contain word  $j$ . Save the result to array  $df$  ( $df$  has shape of  $(n,)$ ).
  - divides word frequency  $x_{i,j}$  by the total number of words in document  $i$ . Save the result as an array named  $tf$  ( $tf$  has shape of  $(m, n)$ ).
  - for each  $x_{i,j}$ , calculates  $tf\_idf_{i,j} = \frac{tf_{i,j}}{df_j}$ . The reason is, if a word appears in most documents, it does not have the discriminative power and often is called a stop word. The inverse of  $df$  can downgrade the weight of such words.  $tf\_idf$  has shape of  $(m, n)$
  - Now, please print the following:
    - print the index of the longest document
    - print the indexes of words with the top 3 largest  $df$  values
    - for the longest document, print the indexes of words with top 3 largest values in the  $tf\_idf$  array.
  - return the  $tf\_idf$  array.
- Note, for all the steps, **do not use any loop**. Just use array functions and broadcasting for high performance computation.

```
In [3]: import numpy as np
import pandas as pd
```

```
In [50]: def analyze_tf(X):  
    tf_idf = None  
  
    # Add your code here  
  
    #print index of the Longest document  
  
    #print indexes of words with the top 3 largest df values  
  
    #return index of top_3 words in the Longest document  
  
    return tf_idf
```

```
In [51]: # dtm.csv is a csv file for test.  
         # It contains word counts in a few documents
```

```
dtm = pd.read_csv("dtm.csv")
print(dtm.head())
analyze_tf(dtm.values)
```

	texas	freeze	leaves	millions	north	mexico	without	power	frozen	\
0	1	1.0	1.0	1.0	1.0	1.0	1.0	1	0.0	
1	1	0.0	0.0	0.0	0.0	0.0	0.0	1	1.0	
2	1	0.0	0.0	0.0	0.0	0.0	1.0	1	0.0	

	wind	turbines	contribute	rolling	blackouts	across	2.7	million \
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
1	1.0	1.0	1.0	1.0	1.0	1.0	0.0	0.0
2	0.0	0.0	0.0	0.0	0.0	0.0	1.0	1.0

```

      people  winter  storm
0      0.0      0.0      0.0
1      0.0      0.0      0.0
2      1.0      1.0      1.0

```

Indexes of the longest documents: 1

Indexes of words with the top 3 largest df values: [0 7 6]

```
Indexes of words with top 3 largest tf_idf values in the longest document: [ 9 11  8]
```

```
Out[51]: array([[0.04166667, 0.125      , 0.125      , 0.125      , 0.125      ,
0.125      , 0.0625     , 0.04166667, 0.      , 0.      ,
0.      , 0.      , 0.      , 0.      , 0.      ,
0.      , 0.      , 0.      , 0.      , 0.      ],
[0.03703704, 0.      , 0.      , 0.      , 0.      ,
0.      , 0.      , 0.03703704, 0.11111111, 0.11111111,
0.11111111, 0.11111111, 0.11111111, 0.11111111, 0.11111111,
0.      , 0.      , 0.      , 0.      , 0.      ],
[0.04166667, 0.      , 0.      , 0.      , 0.      ,
0.      , 0.0625     , 0.04166667, 0.      , 0.      ,
0.      , 0.      , 0.      , 0.      , 0.      ,
0.125      , 0.125      , 0.125      , 0.125      , 0.125      ]])
```

## 2. Define a function to analyze car dataset using pandas

- Define a function named `emotion_analysis` to do the follows:
  - Read "emotion.csv" as a dataframe with the first row in the csv file as column names
  - Count the number of samples labeled for each emotion (i.e. each value in the column "emotion"). Print the counts.
  - Create a new column called `length` to store the number of words in the text column. (hint: "apply" function to split the text by space and then count elements in the resulting list)

- Show the min, max, and mean values of sadness, happiness, and text length for each emotion. Print the results
- Create a cross tabulation to show the average anxiety score of each emotion and each worry value. Use "emotion" as row index and "worry" as column index. Print the table.
  - This function does not have any return. Just print out the result of each calculation step.

```
In [52]: def emotion_analysis():
```

```
# add your code
```

```
In [20]: emotion_analysis()
```

```
===The number of samples labeled for each emotion===
```

```
Anxiety    1381
Sadness     357
Relaxation  333
Fear        230
Anger       107
Happiness   39
Desire      27
Disgust     17
Name: emotion, dtype: int64
```

```
=== min, max, and mean values of sadness, happiness, and text length for each emotion===
```

emotion	sadness			happiness			length		
	mean	amin	amax	mean	amin	amax	mean	amin	amax
Anger	5.672897	1	9	3.177570	1	8	122.084112	88	374
Anxiety	5.719768	1	9	3.333816	1	9	118.066618	59	541
Desire	4.148148	1	8	4.925926	2	8	150.259259	89	1018
Disgust	4.764706	1	8	3.764706	1	6	108.411765	58	158
Fear	6.565217	1	9	3.056522	1	9	118.852174	80	322
Happiness	2.666667	1	9	7.230769	4	9	122.461538	92	272
Relaxation	2.858859	1	9	5.369369	1	9	119.696697	1	292
Sadness	7.436975	2	9	3.112045	1	9	122.117647	85	544

```
=== Cross tabulation of anxiety score by emotion and worry ===
```

worry	1	2	3	4	5	6	7 \
emotion							
Anger	5.50	1.250000	3.416667	4.222222	5.125000	5.833333	6.040000
Anxiety	3.00	7.000000	5.222222	5.941176	6.149425	6.938931	7.597166
Desire	NaN	1.000000	3.142857	3.250000	3.666667	6.375000	6.000000
Disgust	NaN	5.000000	NaN	4.500000	5.000000	4.800000	5.750000
Fear	NaN	NaN	6.000000	5.000000	5.714286	5.454545	6.918033
Happiness	1.00	3.000000	2.400000	3.000000	1.000000	2.625000	4.000000
Relaxation	1.25	1.777778	2.301887	3.029851	3.090909	4.137931	4.219512
Sadness	2.50	3.857143	3.368421	4.606061	4.882353	5.325301	6.156863

worry	8	9
emotion		
Anger	7.000000	7.437500
Anxiety	8.220126	8.770186
Desire	NaN	NaN
Disgust	NaN	NaN
Fear	7.643836	8.256757
Happiness	2.000000	6.500000
Relaxation	3.562500	4.200000
Sadness	6.565217	7.258065

### 3 (Bonus). Calculate Word Cooccurrences

A word cooccurrence is defined as two words, say  $w_1, w_2$ , both appear in a document (but they are not necessarily next to each other). If two words frequently appear together in a list of documents, for example, "sport" and "game", these words can become topic words.

Define a function, `find_coocur`, as follows:

- take an array input similar to that of Q1, e.g.  $X$  of shape  $(m, n)$
- calculate cooccurrences between any pair of words,  $w_1, w_2$  in the  $m$  documents, where  $w_1 \neq w_2$ . Save the coocurrences as an array of shape  $(n, n)$
- return the cooccurrence array

Again, **do not use any loop**. Just use array functions and broadcasting for high performance computation.

```
In [37]: def find_coocur(x):
```

```
c = None
```

```
# add your code
```

```
return c
```

```
In [38]: # x is a test array
x = np.array([[1,0,2,0], [1,1,0,1], [2,0,1,1]])
print(x)
```

```
# For this toy example, w_0 and w_3 coocur 1 time in d_1,
# and 2 times d_2. Therefore, the total cooccurrences is 3.
```

```
find_coocur(x)
```

```
[[1 0 2 0]
 [1 1 0 1]
 [2 0 1 1]]
```

```
Out[38]: array([[0, 1, 4, 3],
               [1, 0, 0, 1],
               [4, 0, 0, 1],
               [3, 1, 1, 0]])
```

```
In [42]: # Structure of your solution to Assignment 1
```

```
import numpy as np
import pandas as pd

# best practice to test your class
# if your script is exported as a module,
# the following part is ignored
# this is equivalent to main() in Java

if __name__ == "__main__":

    # Test Question 1
    print("\n")
    print("=== Test Question 1 ===")

    dtm = pd.read_csv("dtm.csv")
    x = dtm.values
    analyze_tf(x)

    print("\n")
    print("=== Test Question 2 ===")
    emotion_analysis()

    print("\n")
    print("=== Test Question 3 ===")
    print(find_coocur(x))
```

```
=== Test Question 1 ===
Indexes of longest three documents: [4 3 2]
Indexes of words with the top 3 largest df values: [0 1 4]
Indexes of words with top 3 largest values: [[ 2  3  0]
 [ 6  7  8]
 [21 19 18]]
```

```
=== Test Question 2 ===
===The number of samples labeled for each emotion===
Anxiety      1381
Sadness       357
Relaxation    333
Fear          230
Anger         107
Happiness      39
Desire         27
Disgust        17
Name: emotion, dtype: int64
```

```
=== min, max, and mean values of sadness, happiness, and text length for each emotion===
```

emotion	sadness			happiness			length		
	mean	amin	amax	mean	amin	amax	mean	amin	amax
Anger	5.672897	1	9	3.177570	1	8	122.084112	88	374
Anxiety	5.719768	1	9	3.333816	1	9	118.066618	59	541
Desire	4.148148	1	8	4.925926	2	8	150.259259	89	1018
Disgust	4.764706	1	8	3.764706	1	6	108.411765	58	158
Fear	6.565217	1	9	3.056522	1	9	118.852174	80	322
Happiness	2.666667	1	9	7.230769	4	9	122.461538	92	272
Relaxation	2.858859	1	9	5.369369	1	9	119.696697	1	292
Sadness	7.436975	2	9	3.112045	1	9	122.117647	85	544

```
=== Cross tabulation of anxiety score by emotion and worry ===
```

worry	1	2	3	4	5	6	7	\
emotion								
Anger	5.50	1.250000	3.416667	4.222222	5.125000	5.833333	6.040000	
Anxiety	3.00	7.000000	5.222222	5.941176	6.149425	6.938931	7.597166	
Desire	NaN	1.000000	3.142857	3.250000	3.666667	6.375000	6.000000	
Disgust	NaN	5.000000	NaN	4.500000	5.000000	4.800000	5.750000	
Fear	NaN	NaN	6.000000	5.000000	5.714286	5.454545	6.918033	
Happiness	1.00	3.000000	2.400000	3.000000	1.000000	2.625000	4.000000	
Relaxation	1.25	1.777778	2.301887	3.029851	3.090909	4.137931	4.219512	
Sadness	2.50	3.857143	3.368421	4.606061	4.882353	5.325301	6.156863	

```
worry      8      9
emotion
Anger      7.000000  7.437500
Anxiety    8.220126  8.770186
Desire      NaN      NaN
Disgust     NaN      NaN
Fear       7.643836  8.256757
Happiness  2.000000  6.500000
Relaxation  3.562500  4.200000
Sadness    6.565217  7.258065
```

```
=== Test Question 3 ===
[[0. 2. 1. 1. 1. 1. 0. 0. 0. 0. 0. 0. 1. 1. 1. 0. 0. 0. 0.]
 [2. 0. 1. 1. 1. 1. 0. 0. 0. 0. 0. 0. 0. 1. 1. 1. 0. 0. 0.]
 [1. 1. 0. 1. 1. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
```

In [ ]: